



Insurance companies invest a lot of time and money into optimizing their pricing and accurately estimating the likelihood that customers will make a claim. In many countries insurance it is a legal requirement to have car insurance in order to drive a vehicle on public roads, so the market is very large!

(Source: https://www.accenture.com/_acnmedia/pdf-84/accenture-machine-learning-insurance.pdf)

Knowing all of this, On the Road car insurance have requested your services in building a model to predict whether a customer will make a claim on their insurance during the policy period. As they have very little expertise and infrastructure for deploying and monitoring machine learning models, they've asked you to identify the single feature that results in the best performing model, as measured by accuracy, so they can start with a simple model in production.

They have supplied you with their customer data as a csv file called `car_insurance.csv`, along with a table detailing the column names and descriptions below.

The dataset

Column	Description
<code>id</code>	Unique client identifier
<code>age</code>	<p>Client's age:</p> <ul style="list-style-type: none"> 0 : 16-25 1 : 26-39 2 : 40-64 3 : 65+
<code>gender</code>	<p>Client's gender:</p> <ul style="list-style-type: none"> 0 : Female 1 : Male
<code>driving_experience</code>	<p>Years the client has been driving:</p> <ul style="list-style-type: none"> 0 : 0-9 1 : 10-19 2 : 20-29 3 : 30+
<code>education</code>	<p>Client's level of education:</p> <ul style="list-style-type: none"> 0 : No education 1 : High school 2 : University
<code>income</code>	<p>Client's income level:</p> <ul style="list-style-type: none"> 0 : Poverty 1 : Working class 2 : Middle class 3 : Upper class
<code>credit_score</code>	Client's credit score (between zero and one)
<code>vehicle_ownership</code>	<p>Client's vehicle ownership status:</p> <ul style="list-style-type: none"> 0 : Does not own their vehicle (paying off finance) 1 : Owns their vehicle
<code>vehicle_year</code>	<p>Year of vehicle registration:</p> <ul style="list-style-type: none"> 0 : Before 2015 1 : 2015 or later
<code>married</code>	<p>Client's marital status:</p> <ul style="list-style-type: none"> 0 : Not married 1 : Married
<code>children</code>	Client's number of children
<code>postal_code</code>	Client's postal code
<code>annual_mileage</code>	Number of miles driven by the client each year
<code>vehicle_type</code>	<p>Type of car:</p> <ul style="list-style-type: none"> 0 : Sedan 1 : Sports car
<code>speeding_violations</code>	Total number of speeding violations received by the client
<code>duis</code>	Number of times the client has been caught driving under the influence of alcohol
<code>past_accidents</code>	Total number of previous accidents the client has been involved in
<code>outcome</code>	Whether the client made a claim on their car insurance (response variable):

How likely are you to recommend DataLab to a friend or co-worker?



```
# Import required modules
import pandas as pd
import numpy as np
from statsmodels.formula.api import logit

# Start coding!
data = pd.read_csv('car_insurance.csv')
data.isna().any()
data['annual_mileage'].isna().sum()
data['credit_score'].isna().sum()
data.head()
```

...	↑↓	...	↑↓	...	↑↓	...	↑↓	driving_experie...	...	↑↓	ed...	...	↑↓	income	...	↑↓	credi...	...	↑↓	vehicle_owne...	...	↑↓	ve
0		569520		3		0		0-9y			high school		upper class		0.6290273139						1	aft	
1		750365		0		1		0-9y			none		poverty		0.357757117						0	be	
2		199901		0		0		0-9y			high school		working class		0.4931457852						1	be	
3		478866		0		1		0-9y			university		working class		0.2060128507						1	be	
4		731664		1		1		10-19y			none		working class		0.3883658882						1	be	

Rows: 5

Expand

```
#driving experience
data['driving_experience'] = data['driving_experience'].astype('category').cat.codes
mdl1 = logit('outcome ~ driving_experience', data = data).fit()
print(mdl1.params)
exp1 = pd.DataFrame({'driving_experience':np.arange(0, 3, 1)})
pre1 = exp1.assign(outcome = mdl1.predict(exp1))
c1 = mdl1.pred_table()
TN1 = c1[0,0]
TP1 = c1[1,1]
FN1 = c1[1,0]
FP1 = c1[0,1]
acc1 = (TN1 + TP1) / (TN1 + TP1 + FN1 + FP1)
print("accuracy: ", acc1)
```

Optimization terminated successfully.
 Current function value: 0.467390
 Iterations 7
 Intercept 0.515887
 driving_experience -1.668485
 dtype: float64
 accuracy: 0.7771

```
#education
data['education'] = data['education'].astype('category').cat.codes
mdl2 = logit('outcome ~ education', data = data).fit()
print(mdl2.params)
exp2 = pd.DataFrame({'education':np.arange(0, 2, 1)})
pre2 = exp2.assign(outcome = mdl2.predict(exp2))
c2 = mdl2.pred_table()
TN2 = c2[0,0]
TP2 = c2[1,1]
FN2 = c2[1,0]
FP2 = c2[0,1]
acc2 = (TN2 + TP2) / (TN2 + TP2 + FN2 + FP2)
print("accuracy: ", acc2)
```

Optimization terminated successfully.
 Current function value: 0.617408
 Iterations 5
 Intercept -0.573689
 education -0.223746
 dtype: float64
 accuracy: 0.6867

How likely are you to recommend DataLab to a friend or co-worker?



```
#income
data['income'] = data['income'].astype('category').cat.codes
mdl3 = logit('outcome ~ income', data = data).fit()
print(mdl3.params)
exp3 = pd.DataFrame({'income':np.arange(0, 3, 1)})
pre3 = exp3.assign(outcome = mdl3.predict(exp3))
c3 = mdl3.pred_table()
TN3 = c3[0,0]
TP3 = c3[1,1]
FN3 = c3[1,0]
FP3 = c3[0,1]
acc3 = (TN3 + TP3) / (TN3 + TP3 + FN3 + FP3)
print("accuracy: ", acc3)
```

Optimization terminated successfully.
 Current function value: 0.620588
 Iterations 5
 Intercept -0.628377
 income -0.101325
 dtype: float64
 accuracy: 0.6867

```
#vehicle_ownership
data['vehicle_ownership'] = data['vehicle_ownership'].astype('category').cat.codes
mdl4 = logit('outcome ~ vehicle_ownership', data = data).fit()
print(mdl4.params)
exp4 = pd.DataFrame({'vehicle_ownership':np.arange(0, 1)})
pre4 = exp4.assign(outcome = mdl4.predict(exp4))
c4 = mdl4.pred_table()
TN4 = c4[0,0]
TP4 = c4[1,1]
FN4 = c4[1,0]
FP4 = c4[0,1]
acc4 = (TN4 + TP4) / (TN4 + TP4 + FN4 + FP4)
print("accuracy: ", acc4)
```

Optimization terminated successfully.
 Current function value: 0.552412
 Iterations 5
 Intercept 0.322231
 vehicle_ownership -1.724745
 dtype: float64
 accuracy: 0.7351

```
#vehicle_year
data['vehicle_year'] = data['vehicle_year'].astype('category').cat.codes
mdl5 = logit('outcome ~ vehicle_year', data = data).fit()
print(mdl5.params)
exp5 = pd.DataFrame({'vehicle_year':np.arange(0, 1)})
pre5 = exp5.assign(outcome = mdl5.predict(exp5))
c5 = mdl5.pred_table()
TN5 = c5[0,0]
TP5 = c5[1,1]
FN5 = c5[1,0]
FP5 = c5[0,1]
acc5 = (TN5 + TP5) / (TN5 + TP5 + FN5 + FP5)
print("accuracy: ", acc5)
```

Optimization terminated successfully.
 Current function value: 0.572668
 Iterations 6
 Intercept -2.127052
 vehicle_year 1.735442
 dtype: float64
 accuracy: 0.6867

How likely are you to recommend DataLab to a friend or co-worker?



```
#credit score
data['credit_score'] = data['credit_score']
mdl6 = logit('outcome ~ credit_score', data = data).fit()
print(mdl6.params)
exp6 = pd.DataFrame({'vehicle_year':np.arange(0, 1)})
pre6 = exp6.assign(outcome = mdl6.predict(exp5))
c5 = mdl6.pred_table()

Optimization terminated successfully.
    Current function value: 0.567469
    Iterations 6
Intercept      1.949519
credit_score   -5.479372
dtype: float64
```

```
NameError                               Traceback (most recent call last)
File /usr/local/lib/python3.8/dist-packages/patsy/compat.py:36, in call_and_wrap_exc(msg, origin, f, *args, **kwargs)
    35     try:
--> 36         return f(*args, **kwargs)
    37     except Exception as e:

File /usr/local/lib/python3.8/dist-packages/patsy/eval.py:165, in EvalEnvironment.eval(self, expr, source_name, inner_namespace)
    164     code = compile(expr, source_name, "eval", self.flags, False)
--> 165     return eval(code, {}, VarLookupDict([inner_namespace]
    166                                         + self._namespaces))

File <string>:1

NameError: name 'credit_score' is not defined
```

The above exception was the direct cause of the following exception:

```
PatsyError                               Traceback (most recent call last)
File /usr/local/lib/python3.8/dist-packages/statsmodels/base/model.py:1104, in Results._transform_predict_exog(self, exog, transform)
    1103     try:
--> 1104         exog = dmatrix(design_info, exog, return_type="dataframe")
    1105     except Exception as exc:

File /usr/local/lib/python3.8/dist-packages/patsy/highlevel.py:290, in dmatrix(formula_like, data, eval_env, NA_action, return_type)
    289     eval_env = EvalEnvironment.capture(eval_env, reference=1)
--> 290     (lhs, rhs) = _do_highlevel_design(formula_like, data, eval_env,
    291                                         NA_action, return_type)
    292     if lhs.shape[1] != 0:

File /usr/local/lib/python3.8/dist-packages/patsy/highlevel.py:167, in _do_highlevel_design(formula_like, data, eval_env, NA_action, return_type)
    166     if design_infos is not None:
--> 167         return build_design_matrices(design_infos, data,
    168                                         NA_action=NA_action,
    169                                         return_type=return_type)
    170     else:
    171         # No builders, but maybe we can still get matrices

File /usr/local/lib/python3.8/dist-packages/patsy/build.py:888, in build_design_matrices(design_infos, data, NA_action, return_type, dtype)
    887     if factor_info not in factor_info_to_values:
--> 888         value, is_NA = _eval_factor(factor_info, data, NA_action)
    889         factor_info_to_isNAs[factor_info] = is_NA

File /usr/local/lib/python3.8/dist-packages/patsy/build.py:63, in _eval_factor(factor_info, data, NA_action)
    62     factor = factor_info.factor
--> 63     result = factor.eval(factor_info.state, data)
    64     # Returns either a 2d ndarray, or a DataFrame, plus is_NA mask
```

How likely are you to recommend DataLab to a friend or co-worker?



```
File /usr/local/lib/python3.8/dist-packages/patsy/eval.py:547, in EvalFactor._eval(self, code, memorize_state, data)
  546     inner_namespace = VarLookupDict([data, memorize_state["transforms"]])
--> 547     return call_and_wrap_exc("Error evaluating factor",
  548                             self,
  549                             memorize_state["eval_env"].eval,
  550                             code,
  551                             inner_namespace=inner_namespace)
```

```
File /usr/local/lib/python3.8/dist-packages/patsy/compat.py:43, in call_and_wrap_exc(msg, origin, f, *args, **kwargs)
   42     # Use 'exec' to hide this syntax from the Python 2 parser:
---> 43     exec("raise new_exc from e")
  44 else:
  45     # In python 2, we just let the original exception escape -- better
  46     # than destroying the traceback. But if it's a PatsyError, we can
  47     # at least set the origin properly.
```

File <string>:1

```
PatsyError: Error evaluating factor: NameError: name 'credit_score' is not defined
outcome ~ credit_score
          ^^^^^^^^^^
```

During handling of the above exception, another exception occurred:

```
PatsyError                                     Traceback (most recent call last)
Cell In[8], line 6
  4     print(mdl6.params)
  5     exp6 = pd.DataFrame({'vehicle_year':np.arange(0, 1)})
----> 6     pre6 = exp6.assign(outcome = mdl6.predict(exp5))
  7     c5 = mdl6.pred_table()
```

```
File /usr/local/lib/python3.8/dist-packages/statsmodels/base/model.py:1173, in Results.predict(self, exog, transform, *args, **kwargs)
1128     def predict(self, exog=None, transform=True, *args, **kwargs):
1129         """
1130         Call self.model.predict with self.params as the first argument.
1131
1132     (...)
1133     returned prediction.
1134     """
1135     exog, exog_index = self._transform_predict_exog(exog,
1136                                                     transform=transform)
1137     predict_results = self.model.predict(self.params, exog, *args,
1138                                         **kwargs)
1139     if exog_index is not None and not hasattr(predict_results,
1140                                                 'predicted_values'):
```

```
File /usr/local/lib/python3.8/dist-packages/statsmodels/base/model.py:1111, in Results._transform_predict_exog(self, exog, transform)
1105     except Exception as exc:
1106         msg = ('predict requires that you use a DataFrame when '
1107               'predicting from a model\nthat was created using the '
1108               'formula api.\n\nThe original error message returned by patsy is:\n'
1109               '{0}'.format(str(str(exc))))
1110
--> 1111     raise exc.__class__(msg)
1112 if orig_exog_len > len(exog) and not is_dict:
1113     if exog_index is None:
```

```
PatsyError: predict requires that you use a DataFrame when predicting from a model
that was created using the formula api.
```

The original error message returned by patsy is:

```
Error evaluating factor: NameError: name 'credit_score' is not defined
outcome ~ credit_score
          ^^^^^^^^^^
```

How likely are you to recommend DataLab to a friend or co-worker?

